

7. Variational Autoencoder

Young-geun Kim

Department of Mathematical Sciences

MTH345: Deep Generative Models: A Statistical Perspective



ULSAN NATIONAL INSTITUTE OF
SCIENCE AND TECHNOLOGY

Lecture Topics

M1. Foundations of Probabilistic Modelling

1. Fundamentals of Probability Theory
2. Random Variables
3. Statistical Estimation and Hypothesis Testing
4. Basics of Probabilistic Modelling


M2. Likelihood Ratio-based Generative Models

5. Autoregressive Model
6. Energy-based Model
7. Variational Autoencoder
8. Generative Adversarial Network

M3. Generative Models with Other Statistical Distance

9. Integral Probability Metric-based Approaches
10. Wasserstein Distance-based Approaches
11. Score-based Approaches

Lecture Topics



- ▶ In the previous chapter, we reviewed energy-based models. Both training and generation require running a Markov chain (with multi-step Gibbs sampling).
- ▶ In this lecture, we will review another generative model, *variational autoencoders* (VAEs). By applying variational inference, they do not require running Markov chains during either training or testing times.

The figure is edited from Goodfellow (2016).

Lecture Topics

- ▶ Recommended reference materials include:
 - ▶ Kingma and Welling (2014), "Auto-encoding variational Bayes," *Proceedings of the International Conference on Learning Representations (ICLR)*.
 - ▶ Kingma et al. (2019), "An Introduction to Variational Autoencoders," *Foundations and Trends® in Machine Learning*.
 - ▶ Goodfellow (2016), "Nips 2016 tutorial: Generative adversarial networks," *arXiv*.
- ▶ It is also encouraged to refer to *Chapter 7. Variational Autoencoders* in Kim (2026). Compared with that reference, this lecture note includes additional background on variational inference and covers roughly 30% of the VAE research addressed in the cited material.

Table of Contents

1. Variational Inference
2. Variational Autoencoder
3. Multi-modal Variational Autoencoder
4. Temporal Variational Autoencoder

Motivation

- ▶ As a preliminary, we briefly review the framework of *variational inference* (Bishop, 2006).
- ▶ Consider a discrete mixture model with a latent variable $Z \in \{1, \dots, K\}$ and an observed variable X . The generative process is a Gaussian mixture model and is assumed to be known:¹


$$Z \sim \text{Multin}(1, (1/K, \dots, 1/K)^T) \iff p(z) = \frac{1}{K}, \quad \text{for } z = 1, \dots, K, \quad (1)$$

$$X | Z = z \sim \mathcal{N}(z, 1) \iff p(x | z) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(x - z)^2}{2}\right). \quad (2)$$

- ▶ Evaluating the marginal likelihood $p(x)$ poses a severe computational challenge, and variational inference provides a prominent framework to circumvent this issue.


¹From now on, following the convention in variational inference, we use p to denote the PDF regardless of the data type.

Motivation



- The $p(x)$ can be expressed as $p(x) = \sum_{k=1}^K p(x | z)p(z) = \sum_{k=1}^K \phi(x - z)/K$.

Motivation



- Most of the values $p(x | z = k)$ are extremely close to zero, except within the narrow interval $k \in \{47, \dots, 53\}$.

Naive Monte Carlo

- ▶ When K is large or when Z is continuous, the analytical or numerical computation of the marginal likelihood $p(x)$ is often intractable.
- ▶ Here, we first consider an intuitive baseline approach.
- ▶ **Naive Monte Carlo**: Draw $z_{\text{naive}}^{(1)}, \dots, z_{\text{naive}}^{(M)} \stackrel{\text{i.i.d.}}{\sim} p(z)$ and calculate:

$$\hat{p}_{\text{naive}}(x) := \frac{1}{M} \sum_{m=1}^M p(x | z_{\text{naive}}^{(m)}). \quad (3)$$

- ▶ For example, when estimating $p(x = 50)$ with $M = 3$, if the drawn samples are $z_{\text{naive}}^{(1)} = 12$, $z_{\text{naive}}^{(2)} = 50$, and $z_{\text{naive}}^{(3)} = 84$, the empirical calculation yields:

$$\hat{p}_{\text{naive}}(x = 50) = \frac{0 + \frac{1}{\sqrt{2\pi}} + 0}{3} = \frac{1}{3\sqrt{2\pi}} \approx 0.13298. \quad (4)$$

Recall that the true marginal likelihood is

$$p(x = 50) = \frac{1}{100} \sum_{k=1}^{100} \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(50-k)^2}{2}\right) \approx 0.01000.$$

Naive Monte Carlo

Naive Monte Carlo

- ▶ The empirical calculation results heavily depend on the samples falling to $\{47, \dots, 53\}$.
- ▶ Nevertheless, the Monte Carlo estimator is unbiased:

$$\mathbb{E}_{Z_{\text{naive}}^{(m)} \stackrel{\text{i.i.d.}}{\sim} p(z)} [\hat{p}_{\text{naive}}(x)] = \quad (5)$$

Naive Monte Carlo

- Its high variability can be expressed as:

$$\begin{aligned}\text{Var}_{Z_{\text{naive}}^{(m)} \stackrel{\text{i.i.d.}}{\sim} p(z)} [\hat{p}_{\text{naive}}(x)] &= \frac{1}{M} \text{Var}_{Z \sim p(z)} [p(x | Z)] \\ &= \frac{1}{M} \left(\mathbb{E}_{Z \sim p(z)} [p(x | Z)^2] - p(x)^2 \right).\end{aligned}\tag{6}$$

- By applying the CLT, as $M \rightarrow \infty$,

$$\begin{aligned}\sqrt{M} (\hat{p}_{\text{naive}}(x) - p(x)) &\xrightarrow{d} \mathcal{N} \left(0, \text{Var}_{Z \sim p(z)} [p(x | Z)] \right) \\ &= \mathcal{N} \left(0, \mathbb{E}_{Z \sim p(z)} [p(x | Z)^2] - p(x)^2 \right).\end{aligned}\tag{7}$$

Importance Sampling

- ▶ One main source of the high variability of the Naive Monte Carlo approach is that the prior $p(z = k)$ does not reflect the relative contribution of $p(x | z = k)$ to the marginal likelihood $p(x)$.
- ▶ Importance Sampling: Draw $z_{\text{is}}^{(1)}, \dots, z_{\text{is}}^{(M)} \stackrel{\text{i.i.d.}}{\sim} q(z)$ and calculate:

$$\hat{p}_{\text{is}}(x) := \sum_{m=1}^M p(x | z_{\text{is}}^{(m)}) \frac{p(z_{\text{is}}^{(m)})}{Mq(z_{\text{is}}^{(m)})}. \quad (8)$$

The $q(z)$ is referred to as the *proposal distribution*.


- ▶ By appropriately choosing $q(z)$ to yield high probability values where the joint density $p(x, z) = p(x | z)p(z)$ is significant, we can expect to reduce the variability of the estimator.

Importance Sampling

- ▶ The importance sampling estimator is unbiased:

$$\mathbb{E}_{Z_{\text{is}}^{(m)} \sim q(z)} [\hat{p}_{\text{is}}(x)] = \quad (9)$$

Importance Sampling



Importance Sampling

Importance Sampling

- ▶ The variance of the importance sampling estimator can be expressed as:

$$\begin{aligned}\text{Var}_{Z_{\text{is}}^{(m)} \text{i.i.d. } q(z)} [\hat{p}_{\text{is}}(x)] &= \frac{1}{M} \text{Var}_{Z \sim q(z)} \left[\frac{p(Z, x)}{q(Z)} \right] \\ &= \frac{1}{M} \left(\mathbb{E}_{Z \sim q(z)} \left[\left(\frac{p(Z, x)}{q(Z)} \right)^2 \right] - p(x)^2 \right).\end{aligned}\tag{10}$$

- ▶ By applying the CLT, we have:

$$\begin{aligned}\sqrt{M} (\hat{p}_{\text{is}}(x) - p(x)) &\xrightarrow{d} \mathcal{N} \left(0, \text{Var}_{Z \sim q(z)} \left[\frac{p(Z, x)}{q(Z)} \right] \right) \\ &= \mathcal{N} \left(0, \mathbb{E}_{Z \sim q(z)} \left[\left(\frac{p(Z, x)}{q(Z)} \right)^2 \right] - p(x)^2 \right).\end{aligned}\tag{11}$$

Importance Sampling

- Recall that:

$$\text{Var}_{Z_{\text{naive}}^{(m)} \stackrel{\text{i.i.d.}}{\sim} p(z)} [\hat{p}_{\text{naive}}(x)] = \frac{1}{M} \text{Var}_{Z \sim p(z)} [p(x | Z)]. \quad (12)$$

- We derived that:

$$\text{Var}_{Z_{\text{is}}^{(m)} \stackrel{\text{i.i.d.}}{\sim} q(z)} [\hat{p}_{\text{is}}(x)] = \frac{1}{M} \text{Var}_{Z \sim q(z)} \left[\frac{p(Z, x)}{q(Z)} \right]. \quad (13)$$

- By Jensen's inequality, we have:

$$\text{Var}_{Z_{\text{is}}^{(m)} \stackrel{\text{i.i.d.}}{\sim} q(z)} [\hat{p}_{\text{is}}(x)] = \frac{1}{M} \left(\mathbb{E}_{Z \sim q(z)} \left[\left(\frac{p(Z, x)}{q(Z)} \right)^2 \right] - p(x)^2 \right) \geq 0. \quad (14)$$

The equality holds if and only if $q^*(z) = p(z | x)$.

Variational Inference

- ▶ The ideal choice for $q(z)$, the true posterior $q^*(z) = p(z | x)$, is often computationally intractable.
- ▶ **Variational Inference**: We can reformulate the intractable inference problem into a deterministic optimization framework:

$$q_{\phi^*} = \arg \min_{q_{\phi}(z) \in \mathcal{Q}} \mathcal{D}_{\text{KL}}(q_{\phi}(z) \parallel p(z | x)), \quad (15)$$

where $\mathcal{Q}$ denotes a tractable family of distributions, a paradigm also known as *Variational Bayes*.

- ▶ Since q_{ϕ^*} depends on the observed data x , we map each observation to its corresponding data-dependent proposal:


$$x \mapsto q_{\phi^*}(z | x) := \arg \min_{q_{\phi}(z|x) \in \mathcal{Q}} \mathcal{D}_{\text{KL}}(q_{\phi}(z | x) \parallel p(z | x)). \quad (16)$$

- ▶ The *Variational Autoencoder* (VAE) operationalizes this generalized encoding process by parameterizing the mapping via a neural network.

Table of Contents

1. Variational Inference
2. Variational Autoencoder
3. Multi-modal Variational Autoencoder
4. Temporal Variational Autoencoder

Generation with Nonlinear Decoder



- ▶ How about generating data using a decoder? That is, the density is modeled as

$$p_{\theta}(\mathbf{z}, \mathbf{x}) := p(\mathbf{z})p_{\theta}(\mathbf{x} | \mathbf{z}), \quad (17)$$

where $p(\mathbf{z})$ be prior (e.g., Gaussian) and $p_{\theta}(\mathbf{x}|\mathbf{z}) = \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_{\theta}(\mathbf{z}), \boldsymbol{\Sigma}_{\theta}(\mathbf{z}))$ with $\boldsymbol{\mu}_{\theta}(\cdot)$ and $\boldsymbol{\Sigma}_{\theta}(\cdot)$ being neural networks parameterized by θ .

- ▶ This approach extends the decoder ($\boldsymbol{\mu}_{\theta}$) to a probabilistic model. When we sample $\mathbf{Z}$ multiple times from $p(\mathbf{z})$ and sample $\mathbf{X}$ from this (conditional) Gaussian, the distribution of generated $\mathbf{X}$ would be: $p_{\theta}(\mathbf{x}) = \int p(\mathbf{z})p_{\theta}(\mathbf{x}|\mathbf{z})d\mathbf{z}$.

Images are edited from Kim et al. (2021).

Generation with Nonlinear Decoder

- ▶ Due to the flexibility of nonlinear decoder networks, finding the MLE becomes highly challenging:


$$\theta^* = \arg \max_{\theta} \sum_{i=1}^n \log p_{\theta}(\mathbf{x}_i). \quad (18)$$

- ▶ Note that the marginal likelihood $p_{\theta}(\mathbf{x}_i) := \int p(\mathbf{z})p_{\theta}(\mathbf{x}_i | \mathbf{z})d\mathbf{z}$ requires integrating highly nonlinear functions, e.g.,

$$\begin{aligned} & p(\mathbf{z})p_{\theta}(\mathbf{x}_i | \mathbf{z}) \\ & \propto \exp(-\mathbf{z}^T \mathbf{z} / 2) \cdot |\Sigma_{\theta}(\mathbf{z})|^{-1/2} \exp(-(\mathbf{x}_i - \boldsymbol{\mu}_{\theta}(\mathbf{z}))^T \Sigma_{\theta}(\mathbf{z})^{-1} (\mathbf{x}_i - \boldsymbol{\mu}_{\theta}(\mathbf{z})) / 2). \end{aligned} \quad (19)$$

- ▶ Furthermore, the inference with $p_{\theta}(\mathbf{z} | \mathbf{x}) := p_{\theta}(\mathbf{z}, \mathbf{x}) / p_{\theta}(\mathbf{x})$ is also less straightforward.


Variational Autoencoder




- ▶ Variational Autoencoders (VAEs, Kingma and Welling, 2014) utilize variational inference to overcome the aforementioned challenges.
- ▶ They model the data generation process using decoder networks with a diagonal covariance $\Sigma_\theta := \text{diag}(\sigma_1^2, \dots, \sigma_d^2)$:

$$\begin{aligned}
 p_\theta(\mathbf{z}, \mathbf{x}) &= \left(\prod_{j=1}^r p(z_j) \right) \left(\prod_{j=1}^d p_\theta(x_j | \mathbf{z}) \right) \\
 &= \prod_{j=1}^r \left(\frac{1}{\sqrt{2\pi}} \exp \left(-\frac{z_j^2}{2} \right) \right) \prod_{j=1}^d \left(\frac{1}{\sqrt{2\pi\sigma_j^2}} \exp \left(-\frac{(x_j - (\mu_\theta(\mathbf{z}))_j)^2}{2\sigma_j^2} \right) \right).
 \end{aligned}
 \tag{20}$$

Variational Autoencoder



Variational Autoencoder



- ▶ VAEs introduce an encoder, or *inference network*, $q_{\phi}(\mathbf{z}|\mathbf{x})$ to approximate $p_{\theta}(\mathbf{z}|\mathbf{x})$.²
- ▶ By applying the variational inference, we have:

$$\log p_{\theta}(\mathbf{x}) = \log \mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} \left[\frac{p_{\theta}(\mathbf{z}, \mathbf{x})}{q_{\phi}(\mathbf{z} | \mathbf{x})} \right]. \quad (21)$$

²VAEs use $q_{\phi}(\mathbf{z}|\mathbf{x})$, a single function (of $\mathbf{x}$), to approximate the posterior distribution for each new data point, referred to as *amortized variational inference*.

Variational Autoencoder

- ▶ The process of finding the optimal proposal q_{ϕ^*} can be interpreted as maximizing a lower bound of the evidence $\log p_{\theta}(\mathbf{x})$:

$$\log p_{\theta}(\mathbf{x}) \geq \log p_{\theta}(\mathbf{x}) - \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z} | \mathbf{x}) \parallel p_{\theta}(\mathbf{z} | \mathbf{x})). \quad (22)$$

- ▶ Here, the RHS

$$\text{ELBO}(\theta, \phi; \mathbf{x}) := \log p_{\theta}(\mathbf{x}) - \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z} | \mathbf{x}) \parallel p_{\theta}(\mathbf{z} | \mathbf{x})), \quad (23)$$

is referred to as the *evidence lower bound (ELBO)*, which exhibits a computationally tractable form:

$$\text{ELBO}(\theta, \phi; \mathbf{x}) = \mathbb{E}_{\mathbf{Z} \sim q_{\phi}(\mathbf{z} | \mathbf{x})} [\log p_{\theta}(\mathbf{x} | \mathbf{Z})] - \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z} | \mathbf{x}) \parallel p(\mathbf{z})), \quad (24)$$

Variational Autoencoder

► $\text{ELBO}(\theta, \phi; \mathbf{x}) := \log p_{\theta}(\mathbf{x}) - \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z} \mid \mathbf{x}) \parallel p_{\theta}(\mathbf{z} \mid \mathbf{x})).$

Variational Autoencoder

- ▶ In the ELBO,

$$\text{ELBO}(\theta, \phi; \mathbf{x}) = \mathbb{E}_{\mathbf{Z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x} | \mathbf{Z})] - \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z} | \mathbf{x}) \parallel p(\mathbf{z})), \quad (25)$$

the first term on the RHS represents the *reconstruction probability*, and maximizing it can be viewed as training a probabilistic autoencoder.

- ▶ For example, when $q_{\phi}(\mathbf{z} | \mathbf{x}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}_{\phi}(\mathbf{x}), \boldsymbol{\Sigma}_{\phi}(\mathbf{x}))$ and $p_{\theta}(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_{\theta}(\mathbf{z}), \sigma^2 I_d)$, the reconstruction term simplifies to:

$$\mathbb{E}_{\mathbf{Z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x} | \mathbf{Z})] = -\frac{1}{2\sigma^2} \mathbb{E}_{\mathbf{Z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} [\|\mathbf{x} - \boldsymbol{\mu}_{\theta}(\mathbf{Z})\|^2] - \frac{d}{2} \log(2\pi\sigma^2). \quad (26)$$

- ▶ From this perspective, the second term, *KL-penalty term*, acts as a regularizer in the training of autoencoders.

Variational Autoencoder

- ▶ The first term can be approximated via the Monte Carlo method:

$$\int (\log p_{\theta}(\mathbf{x}|\mathbf{z})) q_{\phi}(\mathbf{z}|\mathbf{x}) d\mathbf{z} \approx M^{-1} \sum_{m=1}^M \log p_{\theta}(\mathbf{x}|\mathbf{z}^{(m)}) \quad (27)$$

where $\mathbf{z}^{(m)} \stackrel{\text{i.i.d.}}{\sim} q_{\phi}(\mathbf{z}|\mathbf{x})$, and

$$\log p_{\theta}(\mathbf{x}|\mathbf{z}^{(m)}) = -\frac{1}{2\sigma^2} \sum_{j=1}^d (x_j - (\mu_{\theta}(\mathbf{z}_m))_j)^2 - \frac{d}{2} \log(2\pi\sigma^2). \quad (28)$$

- ▶ Under some conditions, this formulation reduces to the reconstruction error of vanilla autoencoders.
- ▶ In sampling $\mathbf{z}^{(m)} \sim q_{\phi}(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}_{\phi}(\mathbf{x}), \Sigma_{\phi}(\mathbf{x}))$, *reparameterization trick* is applied: $\mathbf{Z}^{(m)} = \boldsymbol{\mu}_{\phi}(\mathbf{x}) + \Sigma_{\phi}(\mathbf{x})^{1/2} \boldsymbol{\epsilon}^{(m)}$ where $\boldsymbol{\epsilon}^{(m)} \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, I_d)$.

Variational Autoencoder

- For the second term, the KL penalty term $\mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$, we can use the following formula:


$$\begin{aligned} & \mathcal{D}_{\text{KL}}(\mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_1, \Sigma_1) || \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_2, \Sigma_2)) \\ &= \frac{1}{2} \left((\boldsymbol{\mu}_2 - \boldsymbol{\mu}_1)^T \Sigma_2^{-1} (\boldsymbol{\mu}_2 - \boldsymbol{\mu}_1) + \text{tr}(\Sigma_2^{-1} \Sigma_1) - \log \frac{|\Sigma_1|}{|\Sigma_2|} - d \right). \end{aligned} \quad (29)$$

See StatProofBook contributors (2026) for details on the proof.

- When $p(\mathbf{z}) = \mathcal{N}(0, I_d)$ and $\Sigma_{\phi}(\mathbf{x}) = \text{diag}(\sigma_{\phi}^2(\mathbf{x}))$, the KL penalty term simplifies to:

$$\mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z} | \mathbf{x}) || p(\mathbf{z})) = \frac{1}{2} \sum_{j=1}^d \left((\mu_{\phi}(\mathbf{x}))_j^2 + (\sigma_{\phi}^2(\mathbf{x}))_j - \log(\sigma_{\phi}^2(\mathbf{x}))_j - 1 \right). \quad (30)$$

Variational Autoencoder



The figure is from Tolstikhin et al. (2017).

Variational Autoencoder




Images are edited from Kingma and Welling (2014).

Table of Contents

1. Variational Inference
2. Variational Autoencoder
3. Multi-modal Variational Autoencoder
4. Temporal Variational Autoencoder

Motivation




- ▶ Data often have multi-modal structures.
- ▶ Multi-modalities share comprehensive information while each modality also possesses its unique information.

The figure is from Gong et al. (2023)

Linear Method

- ▶ (Traditional) Linear methods are usually based on low-rank matrix compression theory:
 1. Multi-set Canonical Correlation Analysis (Kettenring, 1971)
 2. Joint and Individual Variation Explained (Lock et al., 2013)
 3. Multi-block Partial Least Squares (Chen et al., 2009).
- ▶ Multi-modal data may have complicated nonlinear dependencies given latent factors, and we may need to learn mixing mechanisms to analyze new test data.

Joint Multi-modal VAE




- ▶ Joint Multi-modal VAEs (Suzuki et al., 2016) are one of the earliest VAEs-based method, which learn the structure of bi-modal data.
- ▶ They assume that two modalities are generated from the same representations:

$$p_{\theta}(\mathbf{z}, \mathbf{x}^{(1)}, \mathbf{x}^{(2)}) := p(\mathbf{z})p_{\theta}(\mathbf{x}^{(1)}|\mathbf{z})p_{\theta}(\mathbf{x}^{(2)}|\mathbf{z}) \quad (31)$$

where $\mathbf{x}^{(1)}$ and $\mathbf{x}^{(2)}$ are realizations of the bi-modal data.

Joint Multi-modal VAE



- For the inference, they introduce a joint encoder to approximate the posterior distribution of latent variables given the bi-modal data:


$$q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)}). \quad (32)$$

Joint Multi-modal VAE

- ▶ Given the generation and inference processes, the loss function, the negative ELBO w.r.t. $\log p_{\theta}(\mathbf{x}^{(1)}, \mathbf{x}^{(2)})$ can be expressed as:

$$\begin{aligned} & -\log p_{\theta}(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) + \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)})||p_{\theta}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)})) \\ &= -\int (\log p_{\theta}(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}|\mathbf{z})) q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) d\mathbf{z} + \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)})||p(\mathbf{z})) \\ &= -\int (\log p_{\theta}(\mathbf{x}^{(1)}|\mathbf{z}) + \log p_{\theta}(\mathbf{x}^{(2)}|\mathbf{z})) q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) d\mathbf{z} \\ & \quad + \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)})||p(\mathbf{z})). \end{aligned} \tag{33}$$

Joint Multi-modal VAE



- ▶ *Block-missing*, a situation where some modalities are absent, is a chronic problem in multi-modal data analysis (Song et al., 2024).
- ▶ Joint Multi-modal VAEs introduce modality-specific encoders to alleviate this issue.

Joint Multi-modal VAE

- ▶ Let $q_{\phi^{(1)}}(\mathbf{z}|\mathbf{x}^{(1)})$ and $q_{\phi^{(2)}}(\mathbf{z}|\mathbf{x}^{(2)})$ be modality-specific encoders.
- ▶ The modified objective function can be expressed as:

$$\begin{aligned}
 & - \int (\log p_{\theta}(\mathbf{x}^{(1)}|\mathbf{z}) + \log p_{\theta}(\mathbf{x}^{(2)}|\mathbf{z})) q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) d\mathbf{z} + \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) || p(\mathbf{z})) \\
 & + \alpha (\mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) || q_{\phi^{(1)}}(\mathbf{z}|\mathbf{x}^{(1)})) + \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) || q_{\phi^{(2)}}(\mathbf{z}|\mathbf{x}^{(2)}))).
 \end{aligned}
 \tag{34}$$

- ▶ Note that all the posterior distributions, $q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \mathbf{x}^{(2)})$, $q_{\phi^{(1)}}(\mathbf{z}|\mathbf{x}^{(1)})$, and $q_{\phi^{(2)}}(\mathbf{z}|\mathbf{x}^{(2)})$, are modeled as Gaussians, the new added terms have closed-form expressions.

Joint Multi-modal VAE

Input	Generated attributes
	Male : 0.95 Eyeglasses : -0.99 Young : 0.30 Smiling : -0.97 . .
	Male : 0.22 Eyeglasses : -0.99 Young : 0.87 Smiling : -1.00 . .

► Imputation results using $q_{\phi(1)}(\mathbf{z}|\mathbf{x}^{(1)})$ and $p_{\theta}(\mathbf{x}^{(2)}|\mathbf{z})$.

The figure is from Suzuki et al. (2016).

From Bi-modal to Multi-modal


- What happens when the number of modality M increases? When we directly extend the loss function of joint multi-modal VAEs, we have:

$$\begin{aligned}
 & -\log p_{\theta}(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(M)}) + \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(M)})||p_{\theta}(\mathbf{z}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(M)})) \\
 &= -\int \left(\sum_{m=1}^M \log p_{\theta}(\mathbf{x}^{(m)}|\mathbf{z}) \right) q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(M)}) d\mathbf{z} + \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(M)})||p(\mathbf{z})).
 \end{aligned}
 \tag{35}$$

Q: What would be the joint encoder network size for this approach?

Q: Let's consider adding encoders for each observed modality pair. What is the number of possible missing types?

Multi-modal VAE




- ▶ Multi-modal VAEs (Wu and Goodman, 2018) introduce modality-specific VAEs and aggregate information from each modality using product-of-experts techniques (Hinton, 2002).
- ▶ Their generation process can be expressed as:

$$p_{\theta^{(1:M)}}(\mathbf{z}, \mathbf{x}^{(1:M)}) := p(\mathbf{z}) \prod_{m=1}^M p_{\theta^{(m)}}(\mathbf{x}^{(m)}|\mathbf{z}) \quad (36)$$

where $\theta^{(1:M)} := (\theta^{(1)}, \dots, \theta^{(M)})^T$ and $\mathbf{x}^{(1:M)} := (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(M)})^T$.

Multi-modal VAE



- ▶ Product-of-experts techniques enable us to use modality-specific inference networks to build an inference model incorporating all the observed modalities.

Multi-modal VAE

- ▶ When all modalities are observed, the (variational) inference aims to approximate:

$$p_{\theta(1:M)}(\mathbf{z}|\mathbf{x}^{(1:M)}) \propto \frac{\prod_{m=1}^M p_{\theta(m)}(\mathbf{z}|\mathbf{x}^{(m)})}{p(\mathbf{z})^{M-1}} \quad (37)$$

Multi-modal VAE

- ▶ When each modality-specific encoder perfectly approximates the corresponding posterior, i.e., $q_{\phi(m)}(\mathbf{z}|\mathbf{x}^{(m)}) \approx p_{\theta(m)}(\mathbf{z}|\mathbf{x}^{(m)})$ for all m , we can utilize them to approximate $p_{\theta(1:M)}(\mathbf{z}|\mathbf{x}^{(1:M)})$ using Equation (37).
- ▶ In implementation, multi-modal VAEs model the quotient $\tilde{q}_{\phi(m)}(\mathbf{z}|\mathbf{x}^{(m)}) := q_{\phi(m)}(\mathbf{z}|\mathbf{x}^{(m)})/p(\mathbf{z})$ instead of $q_{\phi(m)}(\mathbf{z}|\mathbf{x}^{(m)})$.
- ▶ With quotient networks, the RHS in Equation (37) can be expressed as $p(\mathbf{z}) \prod_{m=1}^M \tilde{q}_{\phi(m)}(\mathbf{z}|\mathbf{x}^{(m)})$. All the distributions of prior and quotient networks are Gaussian, implying a closed-expression.

Multi-modal VAE: Loss Function

► The loss function consists of three negative ELBOs.

1. $(\mathbf{X}^{(1:M)})$ The negative ELBO w.r.t. $\log p_{\theta(1:M)}(\mathbf{x}^{(1:M)})$:


$$\begin{aligned}
 & -\log p_{\theta(1:M)}(\mathbf{x}^{(1:M)}) + \mathcal{D}_{\text{KL}}(q_{\phi(1:M)}(\mathbf{z}|\mathbf{x}^{(1:M)})||p_{\theta(1:M)}(\mathbf{z}|\mathbf{x}^{(1:M)})) \\
 & = -\int \left(\sum_{m=1}^M \log p_{\theta(m)}(\mathbf{x}^{(m)}|\mathbf{z}) \right) q_{\phi(1:M)}(\mathbf{z}|\mathbf{x}^{(1:M)}) d\mathbf{z} + \mathcal{D}_{\text{KL}}(q_{\phi(1:M)}(\mathbf{z}|\mathbf{x}^{(1:M)})||p(\mathbf{z}))
 \end{aligned} \tag{38}$$

where $q_{\phi(1:M)}(\mathbf{z}|\mathbf{x}^{(1:M)}) \propto p(\mathbf{z}) \prod_{m=1}^M \tilde{q}_{\phi(m)}(\mathbf{z}|\mathbf{x}^{(m)})$.

2. $(\mathbf{X}^{(m)})$ The negative ELBO w.r.t. $\log p_{\theta(m)}(\mathbf{x}^{(m)})$


3. $(\mathbf{X}^{(\pi(1):\pi(m))})$ The negative ELBO w.r.t. $\log p_{\theta(\pi(1):\pi(m))}(\mathbf{x}^{(\pi(1):\pi(m))})$ where $\pi(1), \dots, \pi(m)$ are the first m indices from randomly permuted indices.

Multi-modal VAE



- ▶ For example, for a sample having only $\mathbf{X}^{(1)}$ and $\mathbf{X}^{(3)}$, the loss is the summation of the joint ELBO with respect to $(\mathbf{x}^{(1)}, \mathbf{x}^{(3)})^T$ and the marginal ELBOs with respect to $\mathbf{x}^{(1)}$ and $\mathbf{x}^{(3)}$.

Multi-modal VAE



The figure is from Wu and Goodman (2018).

Table of Contents

1. Variational Inference
2. Variational Autoencoder
3. Multi-modal Variational Autoencoder
4. Temporal Variational Autoencoder


Motivation



- Video frames are popular temporal data, i.e., a time series of image frames at each time step.


Examples in video generation. Images are from Szeto et al. (2019).

Motivation



Images are edited from Babaeizadeh et al. (2017).


Auto-Regressive Model



- ▶ Lotter et al. (2016) proposed a deterministic method that minimizes the MSE between true and predicted future frames.
- ▶ This approach can be viewed as an auto-regressive model.
- ▶ Video Pixel Network (Kalchbrenner et al., 2017) extended image auto-regressive models to videos.

The figure is from Kalchbrenner et al. (2017).

Stochastic Variational Video Prediction




- ▶ Stochastic Variational Video Prediction (SV2P) (Babaeizadeh et al., 2017) is one of the earliest VAE-based video generation methods. It proposes two latent variable models: using (i) time-invariant $\mathbf{Z}$ and (ii) time-variant $\mathbf{Z}^{(t)}$.
- ▶ For (i), the generative model can be expressed as:³

$$p_{\theta}(\mathbf{z}, \mathbf{x}^{(1:T)}) = p(\mathbf{z})p_{\theta}(\mathbf{x}^{(1)}|\mathbf{z})p_{\theta}(\mathbf{x}^{(2)}|\mathbf{z}, \mathbf{x}^{(1)}) \cdots p_{\theta}(\mathbf{x}^{(T)}|\mathbf{z}, \mathbf{x}^{(1:(T-1))}). \quad (39)$$

³The SV2P assumes that first few image frames are available at test time.

Stochastic Variational Video Prediction




- SV2P uses $q_\phi(\mathbf{z}|\mathbf{x}^{(1:T)})$ to minimize the negative ELBO w.r.t. $\log p_\theta(\mathbf{x}^{(1:T)})$:

$$\begin{aligned}
 & -\log p_\theta(\mathbf{x}^{(1:T)}) + \mathcal{D}_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}^{(1:T)})||p_\theta(\mathbf{z}|\mathbf{x}^{(1:T)})) \\
 &= -\int \left(\log p_\theta(\mathbf{x}^{(1:T)}|\mathbf{z}) \right) q_\phi(\mathbf{z}|\mathbf{x}^{(1:T)}) d\mathbf{z} + \mathcal{D}_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}^{(1:T)})||p(\mathbf{z})) \\
 &= -\sum_{t=1}^T \int \left(\log p_\theta(\mathbf{x}^{(t)}|\mathbf{z}, \mathbf{x}^{(1:(t-1))}) \right) q_\phi(\mathbf{z}|\mathbf{x}^{(1:T)}) d\mathbf{z} + \mathcal{D}_{\text{KL}}(q_\phi(\mathbf{z}|\mathbf{x}^{(1:T)})||p(\mathbf{z})).
 \end{aligned}$$

(40)


Stochastic Variational Video Prediction



- $p_{\theta}(\mathbf{x}^{(t)} | \mathbf{z}, \mathbf{x}^{(1:(t-1))})$: The parameter θ is shared across all time steps t .

The figure is edited from Babaeizadeh et al. (2017).


Stochastic Variational Video Prediction



- $q_{\phi}(\mathbf{z}|\mathbf{x}^{(1:T)})$: The inference network uses all the image frames to infer the time-invariant latent variable.

The figure is from Babaeizadeh et al. (2017).


Stochastic Variational Video Prediction



- For the case of (ii) a time-variant latent variable $\mathbf{Z}^{(t)}$, the SV2P model formulates the joint distribution as follows:

$$p_{\theta}(\mathbf{z}^{(1:T)}, \mathbf{x}^{(1:T)}) = \prod_{t=1}^T p(\mathbf{z}^{(t)}) \prod_{t=1}^T p_{\theta}(\mathbf{x}^{(t)} | \mathbf{z}^{(t)}, \mathbf{x}^{(1:(t-1))}). \quad (41)$$

Stochastic Variational Video Prediction




- ▶ The time-variant $\mathbf{Z}^{(t)}$ is inferred using the full sequence of frames $\mathbf{X}^{(1:T)}$.
- ▶ The inference network can be formalized as

$$q_{\phi}(\mathbf{z}^{(1:T)} \mid \mathbf{x}^{(1:T)}) = \prod_{t=1}^T q_{\phi}(\mathbf{z}^{(t)} \mid \mathbf{x}^{(1:T)}). \quad (42)$$

- ▶ Here, the conditional functional form $q_{\phi}(\cdot \mid \mathbf{x}^{(1:T)})$ is assumed to be time-invariant. That is, conditioned on the observed video $\mathbf{X}^{(1:T)}$, the latent variables $\mathbf{Z}^{(1)}, \dots, \mathbf{Z}^{(T)}$ are conditionally i.i.d.

Stochastic Variational Video Prediction




- The loss function is the negative ELBO w.r.t. $\log p_{\theta}(\mathbf{x}^{(1:T)})$:


$$\begin{aligned}
 & -\log p_{\theta}(\mathbf{x}^{(1:T)}) + \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}^{(1:T)}|\mathbf{x}^{(1:T)})||p_{\theta}(\mathbf{z}^{(1:T)}|\mathbf{x}^{(1:T)})) \\
 &= \sum_{t=1}^T \left(-\int (\log p_{\theta}(\mathbf{x}^{(t)}|\mathbf{z}^{(t)}, \mathbf{x}^{(1:(t-1))})) q_{\phi}(\mathbf{z}^{(t)}|\mathbf{x}^{(1:T)}) d\mathbf{z}^{(t)} \right. \\
 & \quad \left. + \mathcal{D}_{\text{KL}}(q_{\phi}(\mathbf{z}^{(t)}|\mathbf{x}^{(1:T)})||p(\mathbf{z}^{(t)})) \right)
 \end{aligned} \tag{43}$$

where $q_{\phi}(\mathbf{z}^{(1:T)}|\mathbf{x}^{(1:T)}) = \prod_{t=1}^T q_{\phi}(\mathbf{z}^{(t)}|\mathbf{x}^{(1:T)})$.

Stochastic Variational Video Prediction




Stochastic Variational Video Prediction



Images are edited from Babaeizadeh et al. (2017). PSNR stands for Peak Signal-to-Noise Ratio.

Lecture Topics



The figure is edited from Goodfellow (2016).

References I

- Babaeizadeh, M., Finn, C., Erhan, D., Campbell, R. H., and Levine, S. (2017). Stochastic variational video prediction. *arXiv preprint arXiv:1710.11252*.
- Bishop, C. M. (2006). Pattern recognition and machine learning. *Springer google schola*, 2:1122–1128.
- Chen, K., Reiman, E. M., Huan, Z., Caselli, R. J., Bandy, D., Ayutyanont, N., and Alexander, G. E. (2009). Linking functional and structural brain images with multivariate network analyses: a novel application of the partial least square method. *Neuroimage*, 47(2):602–610.
- Gong, C., Jing, C., Chen, X., Pun, C. M., Huang, G., Saha, A., Nieuwoudt, M., Li, H.-X., Hu, Y., and Wang, S. (2023). Generative ai for brain image computing and brain network computing: a review. *Frontiers in Neuroscience*, 17:1203104.
- Goodfellow, I. (2016). Nips 2016 tutorial: Generative adversarial networks. *arXiv preprint arXiv:1701.00160*.

References II

- Hinton, G. E. (2002). Training products of experts by minimizing contrastive divergence. *Neural computation*, 14(8):1771–1800.
- Kalchbrenner, N., Oord, A., Simonyan, K., Danihelka, I., Vinyals, O., Graves, A., and Kavukcuoglu, K. (2017). Video pixel networks. In *International Conference on Machine Learning*, pages 1771–1779. PMLR.
- Kettenring, J. R. (1971). Canonical analysis of several sets of variables. *Biometrika*, 58(3):433–451.
- Kim, J.-H., Zhang, Y., Han, K., Wen, Z., Choi, M., and Liu, Z. (2021). Representation learning of resting state fmri with variational autoencoder. *NeuroImage*, 241:118423.
- Kim, Y.-g. (2026). Stt997: Deep generative models.
<https://github.com/kyg0910/-STT997-Deep-Generative-Model>. GitHub repository; Accessed: 2026-06-08.

References III

- Kingma, D. P. and Welling, M. (2014). Auto-encoding variational bayes. In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Kingma, D. P., Welling, M., et al. (2019). An introduction to variational autoencoders. *Foundations and Trends® in Machine Learning*, 12(4):307–392.
- Lock, E. F., Hoadley, K. A., Marron, J. S., and Nobel, A. B. (2013). Joint and individual variation explained (jive) for integrated analysis of multiple data types. *The annals of applied statistics*, 7(1):523.
- Lotter, W., Kreiman, G., and Cox, D. (2016). Deep predictive coding networks for video prediction and unsupervised learning. *arXiv preprint arXiv:1605.08104*.
- Song, S., Lin, Y., and Zhou, Y. (2024). Semi-supervised inference for block-wise missing data without imputation. *Journal of Machine Learning Research*, 25(99):1–36.

References IV

- StatProofBook contributors (2026). Kullback-leibler divergence between multivariate normal distributions. <https://statproofbook.github.io/P/mvn-kl.html>. Accessed: 2026-06-08.
- Suzuki, M., Nakayama, K., and Matsuo, Y. (2016). Joint multimodal learning with deep generative models. *arXiv preprint arXiv:1611.01891*.
- Szeto, R., Sun, X., Lu, K., and Corso, J. J. (2019). A temporally-aware interpolation network for video frame inpainting. *IEEE transactions on pattern analysis and machine intelligence*, 42(5):1053–1068.
- Tolstikhin, I., Bousquet, O., Gelly, S., and Schoelkopf, B. (2017). Wasserstein auto-encoders. *arXiv preprint arXiv:1711.01558*.
- Wu, M. and Goodman, N. (2018). Multimodal generative models for scalable weakly-supervised learning. *Advances in neural information processing systems*, 31.